



*Goal: Maximize mean rewards
given a limited amount of pulls*



**(a) Exploration and exploitation
(Multi-armed bandits)**

*Goal: Optimize mean quality
given a limited set of sessions*



(b) Data-driven QoE optimization